



An Introduction to Reinforcement Learning and Multi-arm Bandits

Explore-Exploit Dilemma

B. Ravindran

Reconfigurable and Intelligent Systems (RISE) Group
Department of Computer Science and Engineering
Indian Institute of Technology Madras



Learning to Control

- Familiar models of machine learning
 - Supervised: Classification, Regression, etc.
 - Unsupervised: Clustering, Frequent patterns, etc.
- How did you learn to cycle?
 - Neither of the above
 - Trial and error!
 - Falling down hurts!



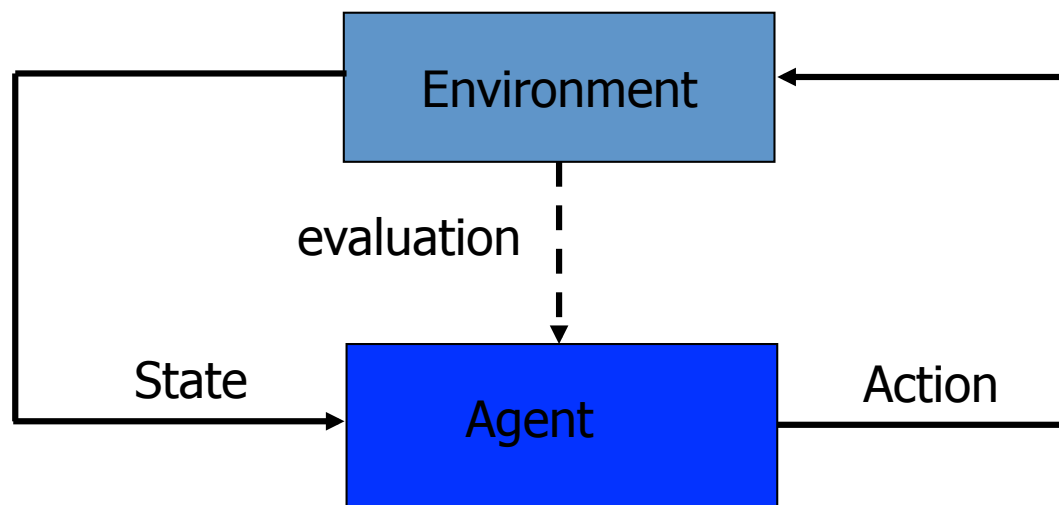


Reinforcement Learning

- A trial-and-error learning paradigm
 - Rewards and Punishments
- Not just an algorithm but a new paradigm in itself
- Learn about a system through interaction
- Inspired by behavioural psychology!




RL Framework



- Learn from close interaction
- Stochastic environment
- Noisy delayed scalar evaluation
- Maximize a measure of long term performance




Not Supervised Learning!



- Very sparse “supervision”
- No target output provided
- No error gradient information available
- Action chooses next state
- Explore to estimate gradient – Trail and error learning




Not Unsupervised Learning



- Sparse “supervision” available
- Pattern detection not primary goal



The Agent-Environment Interface




Agent and environment interact at discrete time steps: $t = 0, 1, 2, \dots$

Agent observes state at step t : $s_t \in S$

produces action at step t : $a_t \in A(s_t)$

gets resulting reward: $r_{t+1} \in \mathfrak{R}$

and resulting next state: s_{t+1}





The Agent Learns a Policy

Policy at step t , π_t :

a mapping from states to action probabilities

$\pi_t(s, a) =$ probability that $a_t = a$ when $s_t = s$

- Reinforcement learning methods specify how the agent changes its policy as a result of experience.
- Roughly, the agent's goal is to get as much reward as it can over the long run.



Goals and Rewards

- Is a scalar reward signal an adequate notion of a goal?—maybe not, but it is surprisingly flexible.
- A goal should specify what we want to achieve, not how we want to achieve it.
- A goal must be outside the agent's direct control—thus outside the agent.
- The agent must be able to measure success:
 - explicitly;
 - frequently during its lifespan.



Returns

Suppose the sequence of rewards after step t is :

$$r_{t+1}, r_{t+2}, r_{t+3}, \dots$$

What do we want to maximize?

In general,

we want to maximize the **expected return**, $E\{R_t\}$, for each step t .

Episodic tasks: interaction breaks naturally into episodes, e.g., plays of a game, trips through a maze.

$$R_t = r_{t+1} + r_{t+2} + \dots + r_T,$$

where T is a final time step at which a **terminal state** is reached, ending an episode.



Returns for Continuing Tasks

Continuing tasks: interaction does not have natural episodes.

Discounted return:


$$R_t = r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots = \sum_{k=0}^{\infty} \gamma^k r_{t+k+1},$$

where $\gamma, 0 \leq \gamma \leq 1$, is the **discount rate**.

shortsighted $0 \leftarrow \gamma \rightarrow 1$ farsighted



An Example



Avoid **failure**: the pole falling beyond a critical angle or the cart hitting end of track.

As an **episodic task** where episode ends upon failure:

reward = +1 for each step before failure

\Rightarrow return = number of steps before failure

As a **continuing task** with discounted return:


reward = -1 upon failure; 0 otherwise

\Rightarrow return = $-\gamma^k$, for k steps before failure

In either case, return is maximized by avoiding failure for as long as possible.



The Markov Property



- “the state” at step t , means whatever information is available to the agent at step t about its environment.
- The state can include immediate “sensations”, highly processed sensations, and structures built up over time from sequences of sensations.
- Ideally, a state should summarize past sensations so as to retain all “essential” information, i.e., it should have the **Markov Property**:

$$\Pr\{s_{t+1} = s', r_{t+1} = r \mid s_t, a_t, r_t, s_{t-1}, a_{t-1}, \dots, r_1, s_0, a_0\} = \Pr\{s_{t+1} = s', r_{t+1} = r \mid s_t, a_t\}$$

for all s', r , and histories $s_t, a_t, r_t, s_{t-1}, a_{t-1}, \dots, r_1, s_0, a_0$.



Markov Decision Processes

- If a reinforcement learning task has the Markov Property, it is basically a Markov Decision Process (MDP).
- If state and action sets are finite, it is a finite MDP.
- To define a finite MDP, you need to give: $M = \langle S, A, P, R \rangle$
 - state and action sets
 - one-step “dynamics” defined by transition probabilities:

$$P_{ss'}^a = \Pr\{s_{t+1} = s' \mid s_t = s, a_t = a\} \quad \text{for all } s, s' \in S, a \in A(s).$$

- reward expectations:

$$R_{ss'}^a = E\{r_{t+1} \mid s_t = s, a_t = a, s_{t+1} = s'\} \quad \text{for all } s, s' \in S, a \in A(s).$$




Markov Decision Processes


- MDP, M , is the tuple: $M = \langle S, A, \Psi, P, R \rangle$
 - S : set of states.
 - A : set of actions.
 - $\Psi \subseteq S \times A$: set of admissible state-action pairs.
 - $P : \Psi \times S \rightarrow [0,1]$: probability of transition.
 - $R : \Psi \rightarrow \mathfrak{R}$: expected reward.
- Policy $\pi : S \rightarrow A$ (can be stochastic)
- Maximize total expected reward.



Example




$$M = \langle S, A, \Psi, P, R \rangle$$






Optimal Policies



$$M = \langle S, A, \Psi, P, R \rangle$$





Solution Methods

- Temporal Difference Methods
 - $TD(\lambda)$
 - Q-learning
 - SARSA
 - Actor-Critic
- Policy Search
 - Policy Gradient Methods
 - Evolutionary algorithms
- Stochastic Dynamic Programming



Applications of RL

- Optimal Control
 - Robot Navigation
 - Helicopters!
 - Chemical Plants
- Combinatorial Optimization
 - Elevator Dispatching
 - VLSI placement and routing
 - Job-shop scheduling
 - Routing algorithms
 - Call admission control
- More
 - Intelligent Tutoring Systems
- Computational Neuroscience
 - Primary mechanism of learning
- Psychology
 - Behavioral and operant conditioning
 - Decision making
- Operations Research
 - Approximate Dynamic Programming
- More
 - Game Playing
 - Dialogue systems

Reinforcement Learning

Lecture 8

Gillian Hayes

1st February 2007



Algorithms for Solving RL: Monte Carlo Methods

- What are they?
- Monte Carlo Policy Evaluation
- First-visit policy evaluation
- Estimating Q-values
- On-policy methods
- Off-policy methods

Monte Carlo Methods

- **Learn** value functions
- **Discover** optimal policies
- Don't require environmental knowledge: $P_{ss'}^a, R_{ss'}^a$,
cf. Dynamic Programming
- Experience : sample sequences of states, actions, rewards s, a, r
: real experience, simulated experience
- Attains optimal behaviour

How Does Monte Carlo Do This?

- Divide experience into episodes
 - all episodes must terminate
e.g. noughts-and-crosses, card games
- Keep estimates of value functions, policies
- Change estimates/policies at end of each episode
 \Rightarrow Keep track of $s_1, a_1, r_1, s_2, a_2, r_2, \dots, s_{T-1}, a_{T-1}, r_{T-1}, s_T$
 $s_T =$ terminating state
- Incremental episode-by-episode
NOT step-by-step cf. DP
- Average **complete** returns – NOT partial returns

Returns

- Return at time t : $R_t = r_{t+1} + r_{t+2} + \dots r_{T-1} + r_T$ for each episode
 r_T is a terminating state
- Average the returns over many episodes starting from some state s .

This gives the value function $V^\pi(s)$ for that state for policy π since the state value $V^\pi(s)$ is the expected cumulative future discounted reward starting in s and following policy π .

Monte Carlo Learning of V^π

MC methods estimate from experience: generate many “plays” from s , observe total reward on each play, average over many plays


1. Initialise

- π = arbitrary policy to be evaluated
- V = arbitrary value function
- $Returns(s)$ an empty list, one for each state s

2. Repeat till values converge

- Generate an episode using π
- For each state appearing in the episode
 - R = return following first occurrence of s
 - Append R to $Returns(s)$
 - $V(s) = \text{average } Returns(s)$

Backup Diagram for MC



One Episode – full episode needed before back-up.
cf DP which backs up after one move
Monte Carlo does **not** bootstrap but
Monte Carlo does sample

- Play many games
- Average returns (first-visit MC) following each state
- \Rightarrow True state-value functions
 - * Easier than DP \Rightarrow That needs $P_{ss'}^a, R_{ss'}^a$
 - * Easier to generate episodes than calculate probabilities

Policy Iteration (Reminder)

- Policy evaluation: Estimate V^π or Q^π for fixed policy π
- Policy improvement: Get a policy better than π

Iterate until optimal policy/value function is reached

So we can do Monte Carlo as the Policy Evaluation step of Policy Iteration because it computes the value function for a given policy. (There are other algorithms we can use.)

First-visit MC vs. Every-visit MC

In each episode observe return following **first** visit to state s

Number of first visits to s must $\rightarrow \infty$

Converges to $V^\pi(s)$

cf. Every-visit MC

Calculate V as the average over return following **every** visit to state s in a set of episodes

Good Properties of MC

Estimates of V for each state are independent

- no bootstrapping

Compute time to calculate changes (i.e. V of each state) is independent of number of states

If values of only a few states needed, generate episodes from these states \Rightarrow can ignore other states

Can learn from actual/simulated experience

Don't need $P_{ss'}^a$, $R_{ss'}^a$,

Estimating Q-Values

$Q^\pi(s, a)$ – similarly to V

Update by averaging returns following first visit to that state-action pair

Problem

If π deterministic, some/many (s, a) never visited

MUST EXPLORE!

So...

- * Exploring starts: start every episode at a different (s, a) pair
- * Or always use ϵ -greedy or ϵ -soft policies
 - stochastic, where $\pi(s, a) > 0$

Optimal Policies – Control Problem

Policy Iteration on Q

$$\pi_0 \rightarrow_{PE} Q^{\pi_0} \rightarrow_{PI} \pi_1 \rightarrow_{PE} Q^{\pi_1} \rightarrow_{PI} \pi_2 \dots \rightarrow_{PI} \pi^* \rightarrow_{PE} Q^*$$

- Policy Improvement: Make π greedy w.r.t. current Q
- Policy Evaluation: As before, with ∞ episodes

Or episode-by-episode iteration. After an episode:

- policy evaluation (back-up)
- improve policy at states in episode
- eventually converges to optimal values and policy

Can use exploring starts: MCES – Monte Carlo Exploring Starts to ensure coverage of state/action space

Algorithm: see e.g. S+B Fig. 5.4

Monte Carlo: Estimating $Q^\pi(s, a)$

- If π deterministic, some (s, a) not visited \Rightarrow can't improve their Q estimates
MUST MAINTAIN EXPLORATION!
- Use exploring starts \rightarrow optimal policy
- Use an ϵ -soft policy
 - ON-POLICY CONTROL \rightarrow ϵ -greedy policy
 - OFF-POLICY CONTROL \rightarrow optimal policy

On-Policy Control

Evaluate and improve the policy used to generate behaviour

Use a soft policy:

$\pi(s, a) > 0 \quad \forall s, \forall a$ GENERAL SOFT POLICY DEFINITION

$\pi(s, a) = \frac{\epsilon}{|A|}$ if a not greedy ϵ -GREEDY

$= 1 - \epsilon + \frac{\epsilon}{|A|}$ if a greedy

$\pi(s, a) \geq \frac{\epsilon}{|A|} \quad \forall s, \forall a$ ϵ -SOFT

POLICY ITERATION

Evaluation: as before *Improvement:* move towards ϵ -greedy policy (not greedy)

Avoids need for exploring starts

ϵ -greedy is “closer” to greedy than other ϵ -soft policies

Off-Policy Control

- Behaviour policy π' generates moves
- But in off-policy control we learn an Estimation policy π . How?

We need to:

- compute the weighted average of returns from behaviour policy
- the weighting factors are the probability of them being in estimation policy,
- i.e. weight each return by relative probability of being generated by π and π'

In detail...

Reinforcement Learning

Lecture 10

Gillian Hayes

8th February 2007



Algorithms for Solving RL: Temporal Difference Learning (TD)

- Incremental Monte Carlo Algorithm
- TD Prediction
- TD vs MC vs DP
- TD for control: SARSA and Q-learning

Incremental Monte Carlo Algorithm

Our first-visit MC algorithm had the steps:

R is the return following our first visit to s

Append R to $Returns(s)$

$V(s) = \text{average}(Returns(s))$

We can implement this incrementally:

$$V(s) = V(s) + \frac{1}{n(s)}[R - V(s)]$$

where $n(s)$ is the number of first visits to s

We can also formulate a constant- α Monte Carlo update:

$$V(s) = V(s) + \alpha[R - V(s)]$$

useful when tracking a non-stationary problem (why?).

Model-Based vs Model-Free Learning

- In RL we're generally trying to learn an optimal policy
- If a model is available, $P_{ss'}^a$, $R_{ss'}^a$, we can calculate optimal policy via dynamic programming
- If no model, either:
 - learn model and then derive optimal policy
(model-based methods) or
 - learn optimal policy without learning model
(model-free methods)
- Temporal difference (TD) learning is a model-free, bootstrapping method based on sampling the state-action space

Temporal Difference Prediction

Policy Evaluation is often referred to as the Prediction Problem: we are trying to predict how much return we'll get from being in state s and following policy π by learning the state-value function V^π .

Monte-Carlo update:

$$V(s_t) \rightarrow V(s_t) + \alpha[R_t - V(s_t)]$$

Target: actual return from s_t to end of episode

Simplest temporal difference update TD(0):

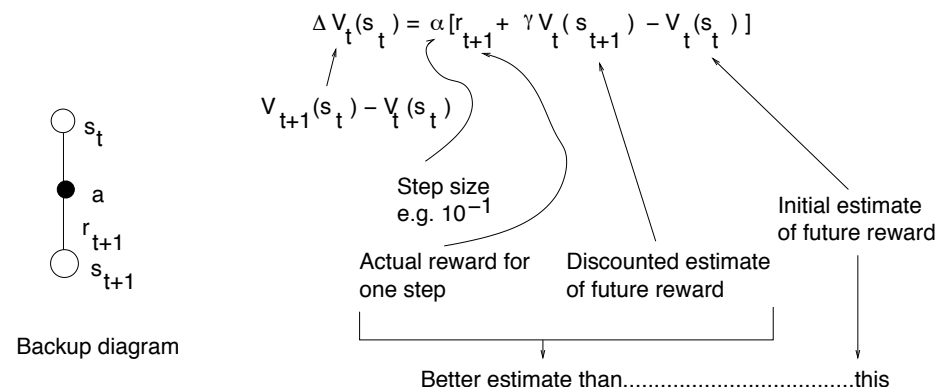
$$V(s_t) \rightarrow V(s_t) + \alpha[r_{t+1} + \gamma V(s_{t+1}) - V(s_t)]$$

Target: estimate of the return

Both have the same form

Temporal Difference Learning

- Doesn't need a model $P_{ss'}^a, R_{ss'}^a$
- Learns directly from experience
- Updates estimates of $V(s)$ based on what happens after visiting state s



TD(0) update:

$$V(s_t) \rightarrow V(s_t) + \alpha[r_{t+1} + \gamma V(s_{t+1}) - V(s_t)]$$

cf Dynamic Programming update:

$$\begin{aligned} V^\pi(s) &= E_\pi\{r_{t+1} + \gamma V^\pi(s_{t+1}) \mid s_t = s\} \\ &= \sum_a \pi(s, a,) \sum_{s'} P_{ss'}^a [R_{ss'}^a + \gamma V^\pi(s')] \end{aligned}$$

Advantages of TD Learning Methods

- Don't need a model of the environment
- On-line and incremental so can be fast
don't need to wait till the end of the episode so need less
memory, computation
- Updates are based on actual experience (r_{t+1})
- Converges to $V^\pi(s)$ – but must decrease step size α as learning continues
- Compare backup diagrams of TD, MC and DP

Bootstrapping, Sampling

TD **bootstraps**: it updates its estimates of V based on other estimates of V

DP also bootstraps

MC does not bootstrap: estimates of complete returns are made at the end of the episode

TD **samples**: its updates are based on one path through the state space

MC also samples

DP does not sample: its updates are based on all actions and all states that can be reached from the updating state

Examples: see e.g. random walk example S+B sect. 6.2

MC vs TD updating: see e.g. S+B sect. 6.3

Difference Between TD and MC Estimates

See S+B Example 6.4:


Suppose you observe the following 8 episodes:

A, 0, B, 0	B, 1
B, 1	B, 1
B, 1	B, 1
B, 1	B, 0

First episode starts in state A, transitions to B getting a reward of 0, and terminates with a reward of 0. Second episode starts in state B and terminates with a reward of 1, etc.

What are the best values for the estimates $V(A)$ and $V(B)$?

Modelling the Underlying Markov Process



$$V(A) = ?$$

TD and MC Estimates


- Batch Monte Carlo (updating after all these episodes are done) gets $V(A) = 0$.
 - This best matches the training data
 - It minimises the mean-square error on the training set
- Consider sequentiality, i.e. A goes to B goes to terminating state; then $V(A) = 0.75$.
 - This is what TD(0) gets
 - Expect that this will produce better estimate of future data even though MC gives the best estimate on the present data

- Is correct for the maximum-likelihood estimate of the model of the Markov process that generates the data, i.e. the best-fit Markov model based on the observed transitions
- Assume this model is correct; estimate the value function – “certainty-equivalence estimate”

TD(0) tends to converge faster because it’s moving towards a “better” estimate.

TD for Control: Learning Q-Values

Learn action values $Q^\pi(s, a)$ for the policy π



SARSA update rule:

$$\Delta Q_t(s_t, a_t) = \alpha[r_{t+1} + \gamma Q_t(s_{t+1}, a_{t+1}) - Q_t(s_t, a_t)]$$

- Choose a behaviour policy π and estimate the Q-values (Q^π) using the SARSA update rule. Change π towards greediness wrt Q^π .
- Use ϵ -greedy or ϵ -soft policies.
- Converges with probability 1 to optimal policy and Q-values if visit all state-action pairs infinitely many times and policy converges to greedy policy, e.g. by arranging for ϵ to tend towards 0.

Remember: learning optimal Q-values is useful since it tells us immediately which is(are) the optimal action(s) – have the highest Q-value

SARSA Algorithm

- Initialise $Q(s, a)$
- Repeat many times
 - Pick s, a
 - Repeat each step to goal
 - * Do a , observe r, s'
 - * Choose a' based on $Q(s', a')$ ϵ -greedy
 - * $Q(s, a) = Q(s, a) + \alpha[r + \gamma Q(s', a') - Q(s, a)]$
 - * $s = s', a = a'$
 - Until s terminal (where $Q(s', a') = 0$)

Use with policy iteration, i.e. change policy each time to be greedy wrt current estimate of Q

Example: windy gridworld, S+B sect. 6.4

Q-Learning

SARSA is an example of **on-policy** learning. Why?

Q-LEARNING is an example of **off-policy** learning

Update rule:


$$\Delta Q_t(s_t, a_t) = \alpha[r_{t+1} + \gamma \max_a Q_t(s_{t+1}, a) - Q_t(s_t, a_t)]$$

Always update using *maximum* Q value available from next state: then $Q \Rightarrow Q^*$, optimal action-value function

Q-Learning Algorithm

- Initialise $Q(s, a)$
- Repeat many times
 - Pick s start state
 - Repeat each step to goal
 - * Choose a based on $Q(s, a)$ ϵ -greedy
 - * Do a , observe r, s'
 - * $Q(s, a) = Q(s, a) + \alpha[r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$
 - * $s = s'$
 - Until s terminal

Backup Diagrams for SARSA and Q-Learning



SARSA backs up using the action a' actually chosen by the behaviour policy.

Q-LEARNING backs up using the Q -value of the action a'^* that is the *best* next action, i.e. the one with the highest Q value, $Q(s', a'^*)$. The action actually chosen by the behaviour policy *and followed* is not necessarily a'^*

Example: The cliff S+B sect. 6.5

Q-Learning vs SARSA

QL: $Q(s, a) = Q(s, a) + \alpha[r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$ off-policy

SARSA: $Q(s, a) = Q(s, a) + \alpha[r + \gamma Q(s', a') - Q(s, a)]$ on-policy

In the cliff-walking task:

QL: learns optimal policy along edge

SARSA: learns a safe non-optimal policy away from edge

ϵ -greedy algorithm

For $\epsilon \neq 0$ **SARSA** performs better online. Why?

For $\epsilon \rightarrow 0$ gradually, both converge to optimal.